




Extreme Cleverness

Functional Data Structures in Scala



Agenda



- Functional data structures
- Implementations
 - Sequential
 - Associative
- Modern computer architecture

Functional Data Struct.

- Immutable, immutable, immutable

Functional Data Struct.

- Immutable, immutable, immutable
- What we want...
 - Comparable asymptotic performance
 - Non-degraded versions
 - (full persistence)

Functional Data Struct.

- Immutable, immutable, immutable
- What we want...
 - Comparable asymptotic performance
 - Non-degraded versions
 - (full persistence)
- Structural sharing


Sequential

- Singly-Linked List
- Banker's Queue
- 2-3 Finger Tree



Singly-Linked List

`List(a, b, c, d)`



Complexity

$O(1)$	$O(\log n)$	$O(n)$
first		append
prepend		concat
		insert
		last
		nth

Anatomy

- A list is *either...*
 - A “cons” cell with a value and a tail
 - An empty list, called “nil”
- These are the only cases!


```
sealed trait List[+A] {  
  def ::[B >: A](b: B): List[B] =  
    new ::(b, this)  
}
```

```
case class ::[+A](hd: A, tail: List[A]) extends List[A]  
case object Nil extends List[Nothing]
```


Structural Sharing

```
val foo = a :: b :: c :: d :: Nil
```

Structural Sharing




Structural Sharing




`val bar = h :: foo.tail`


Structural Sharing




Structural Sharing




Structural Sharing



Structural Sharing




Structural Sharing



Banker's Queue [1]

Queue(a, b, c, d, e)



Motivation

- We want a functional queue
- Linked list is obvious
 - prepend and last are opposing
 - One will be $O(1)$, the other $O(n)$
- Can we have our cake and eat it too?

Complexity

$O(1)$	$O(\log n)$	$O(n)$
append		concat
last		first
prepend		insert
		nth

Complexity

amortized

$O(1)$	$O(\log n)$	$O(n)$
append		concat
last		first
prepend		insert
		nth

Anatomy

- Naïve persistent queue
- Two lazy singly-linked lists
 - Front list (for dequeue)
 - Rear list (for enqueue)
- Periodically reverse rear into the front
- Lazy amortization

Amortization

- *Most* operations are legitimately fast
 - *Few* operations are very slow
- Laziness distributes the work
- Net result: constant factor degradation
 - Translation: the net average is fast
- Also works without laziness!



```
class BankersQueue[+A](fsize: Int, front: Stream[A],  
                        rsize: Int, rear: Stream[A]) {  
  ...  
}
```

```
object BankersQueue {  
  def check[A](q: BankersQueue[A]) = {  
    if (q.rsize <= q.fsize) {  
      q // already valid  
    } else {  
      val fsize2 = q.fsize + q.rsize  
      val front2 = q.front ++ q.rear.reverse  
  
      new BankersQueue(fsize2, front2, 0, Stream())  
    }  
  }  
}
```

```
class BankersQueue[+A](fsize: Int, front: Stream[A],  
                        rsize: Int, rear: Stream[A]) {  
  
  def enqueue[B >: A](b: B) =  
    check(new BankersQueue(fsize, front,  
                           rsize + 1, b #:: rear))  
  
  def dequeue = front match {  
    case hd #:: tail => {  
      val rem = new BankersQueue(fsize - 1, tail,  
                                  rsize, rear)  
      (hd, check(rem))  
    }  
  
    case _ => throw new NoSuchElementException  
  }  
}
```

2-3 Finger Tree [3]

FingerTree('A' to 'W': _*)





Complexity


$O(1)$	$O(\log n)$	$O(n)$
append	insert	concat
first	nth	
last		
prepend		


Anatomy


- Ideal persistent deque
- Digits of length 1, 2, 3 or 4
 - Head and tail
- Branching factor of 2 or 3
- Recursive tree body



































```
sealed trait FingerTree[+A] {  
  ...  
}  
  
case class Single[+A](a: A) extends FingerTree[A] {  
  ...  
}  
  
case class Deep[+A](prefix: Digit[A],  
                    tree: FingerTree[Node[A]],  
                    suffix: Digit[A])  
  extends FingerTree[A] {  
  ...  
}  
  
case object Empty extends FingerTree[Nothing] {  
  ...  
}
```

```
sealed trait Digit[+A]
```

```
case class One[+A](a1: A) extends Digit[A]
```

```
case class Two[+A](a1: A, a2: A) extends Digit[A]
```

```
case class Three[+A](a1: A, a2: A, a3: A) extends Digit[A]
```

```
case class Four[+A](a1: A, a2: A, a3: A, a4: A)  
    extends Digit[A]
```

```
sealed trait Node[+A]
```

```
case class Node2[+A](a1: A, a2: A) extends Node[A]
```

```
case class Node3[+A](a1: A, a2: A, a3: A) extends Node[A]
```

```
case class Deep[+A](...) extends FingerTree[A] {  
  ...  
  def +:[B >: A](b: B) = prefix match {  
    case Four(d, e, f, g) =>  
      Deep(Two(b, d), Node3(e, f, g) +: tree, suffix)  
  
    case _ => Deep(b +: prefix, tree, suffix)  
  }  
  
  def :+[B >: A](b: B) = suffix match {  
    case Four(g, f, e, d) =>  
      Deep(prefix, tree :+ Node3(g, f, e), Two(d, b))  
  
    case _ => Deep(prefix, tree, suffix :+ b)  
  }  
  ...  
}
```


Associative

- Red-Black Tree



Red-Black Tree [2]

RedBlack(1, 2, 3, 4, 5)



Complexity

$O(1)$	$O(\log n)$	$O(n)$
	get	intersect
	insert	union
	update	


Anatomy

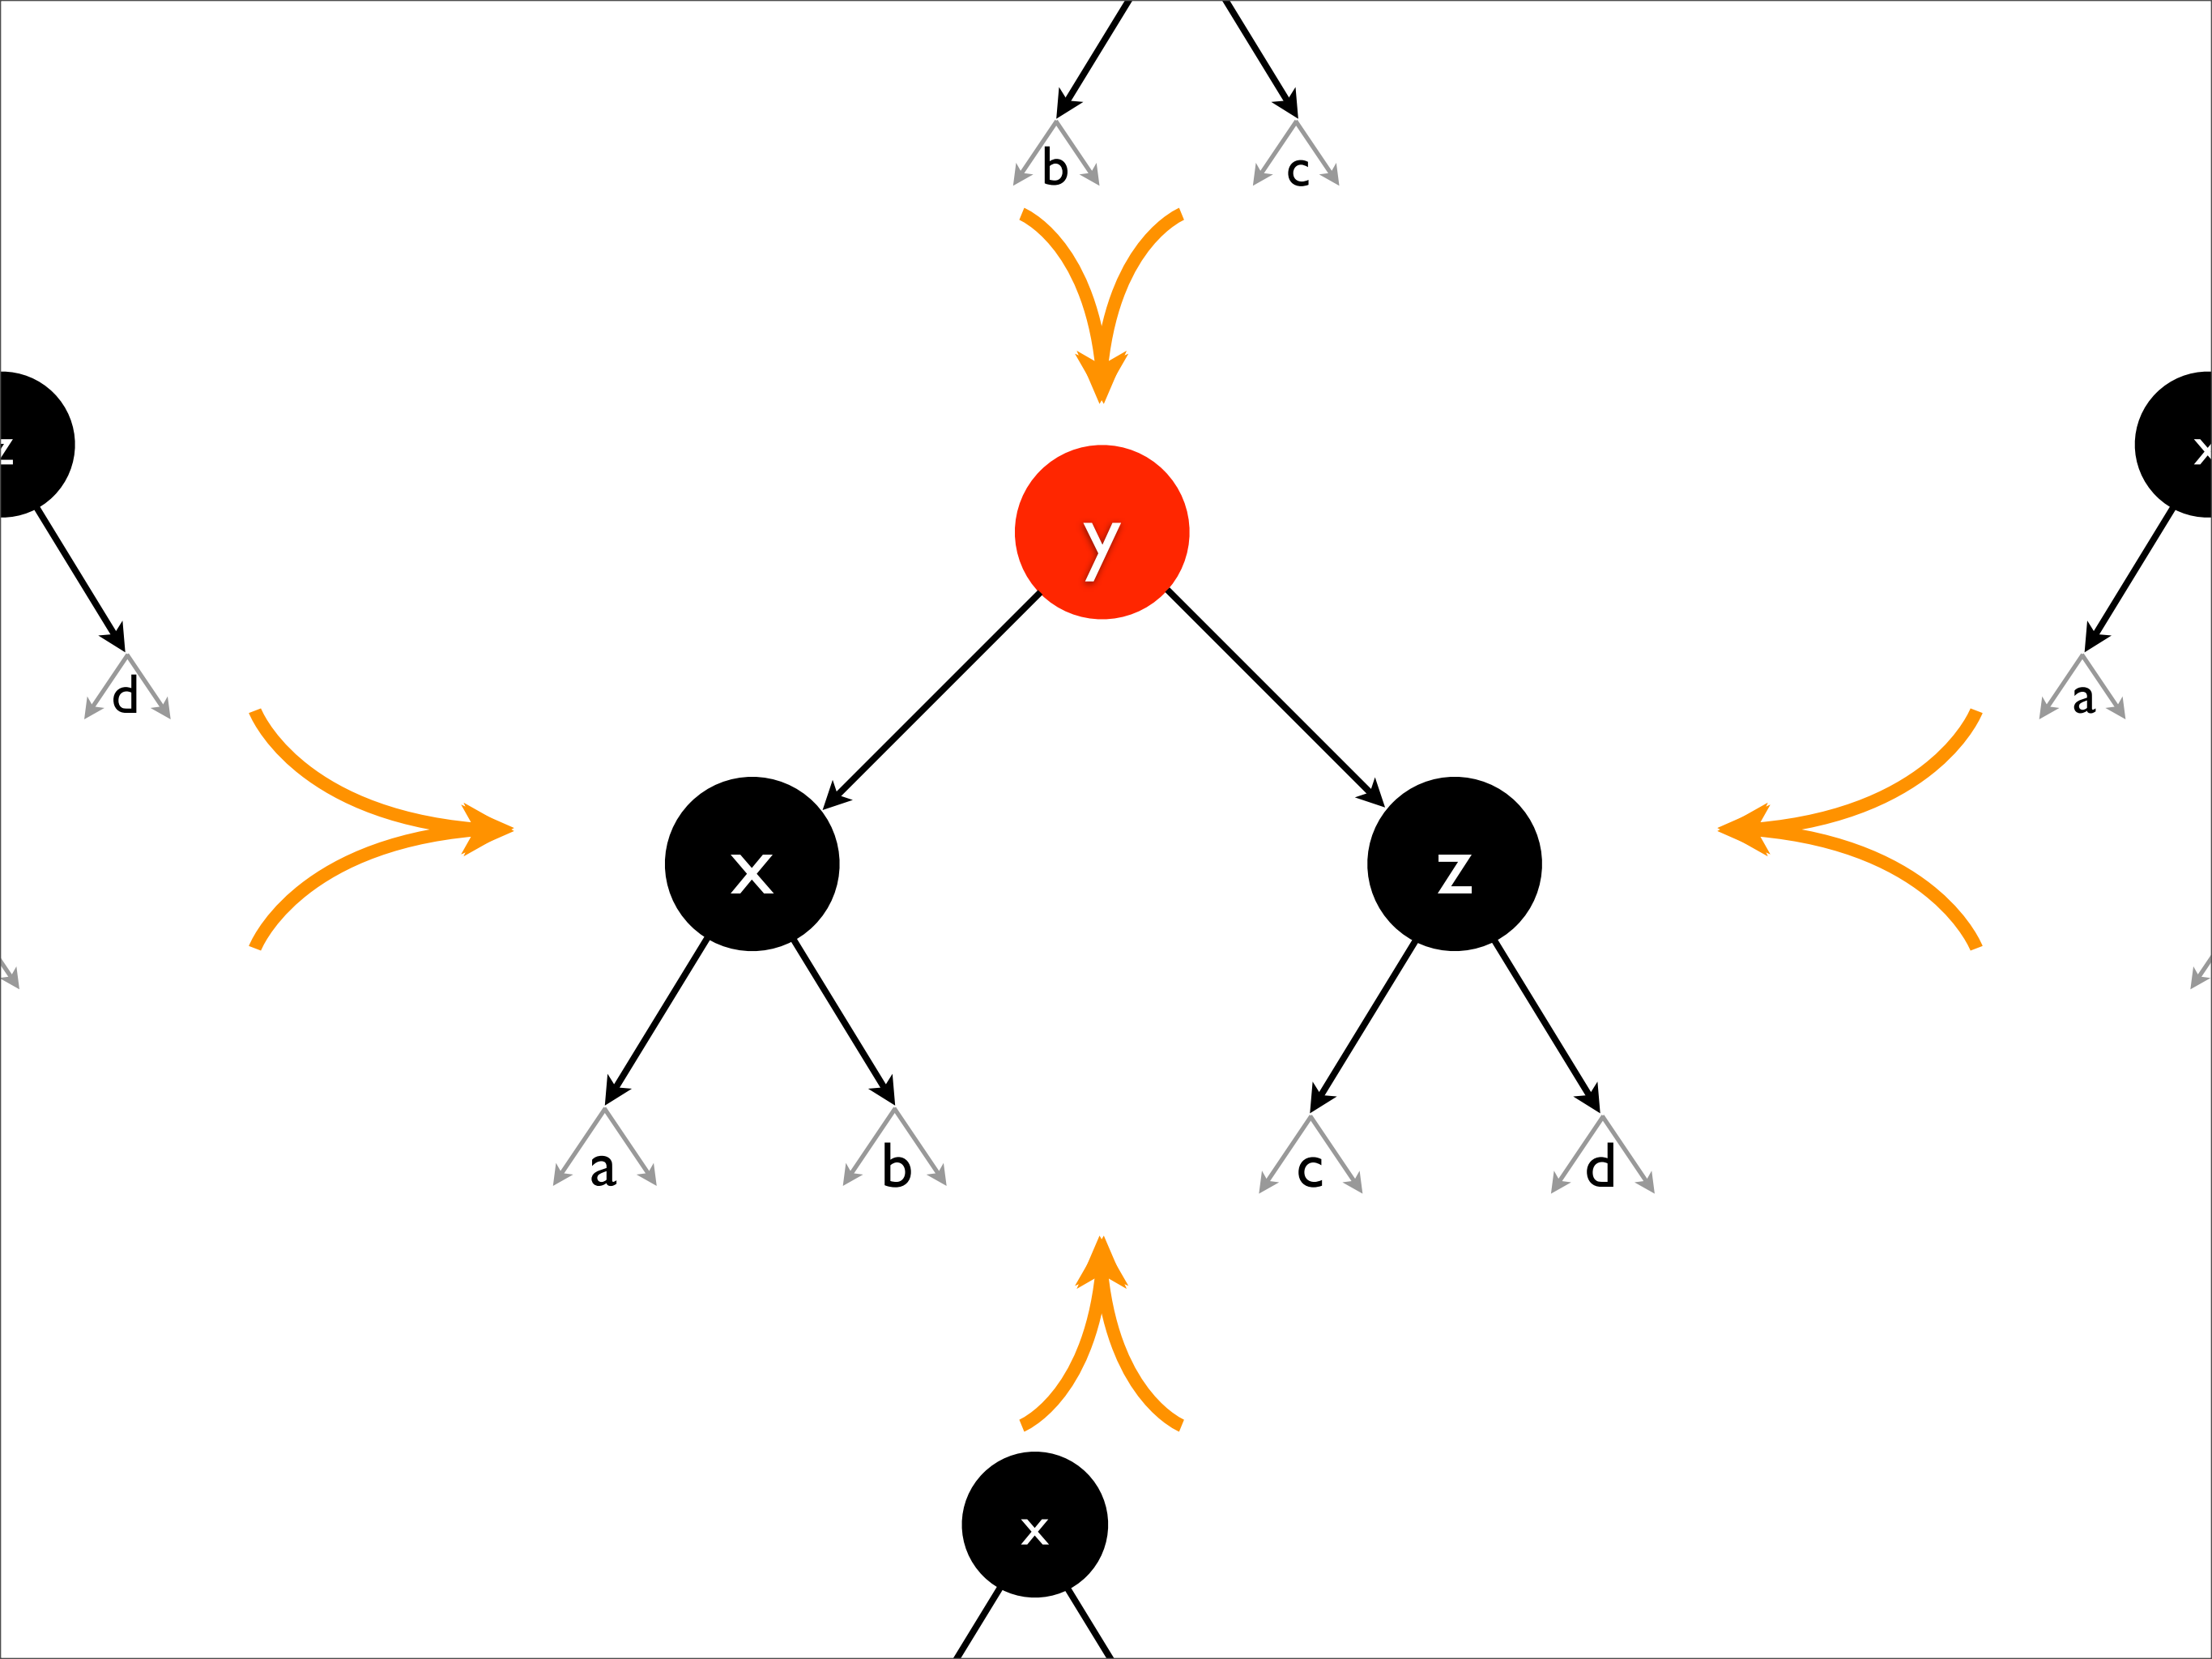
- Balanced binary search tree
- Invariants...
 - Every path from root to a leaf contains the *same* number of **black** nodes
 - No **red** node has a **red** parent
- Need to rebalance after any “update”




```
sealed abstract class Tree[K : Ordering, +V] {  
  val isBlack: Boolean  
  
  def left: Tree[K, V]  
  
  def key: K  
  def value: V  
  
  def right: Tree[K, V]  
}
```

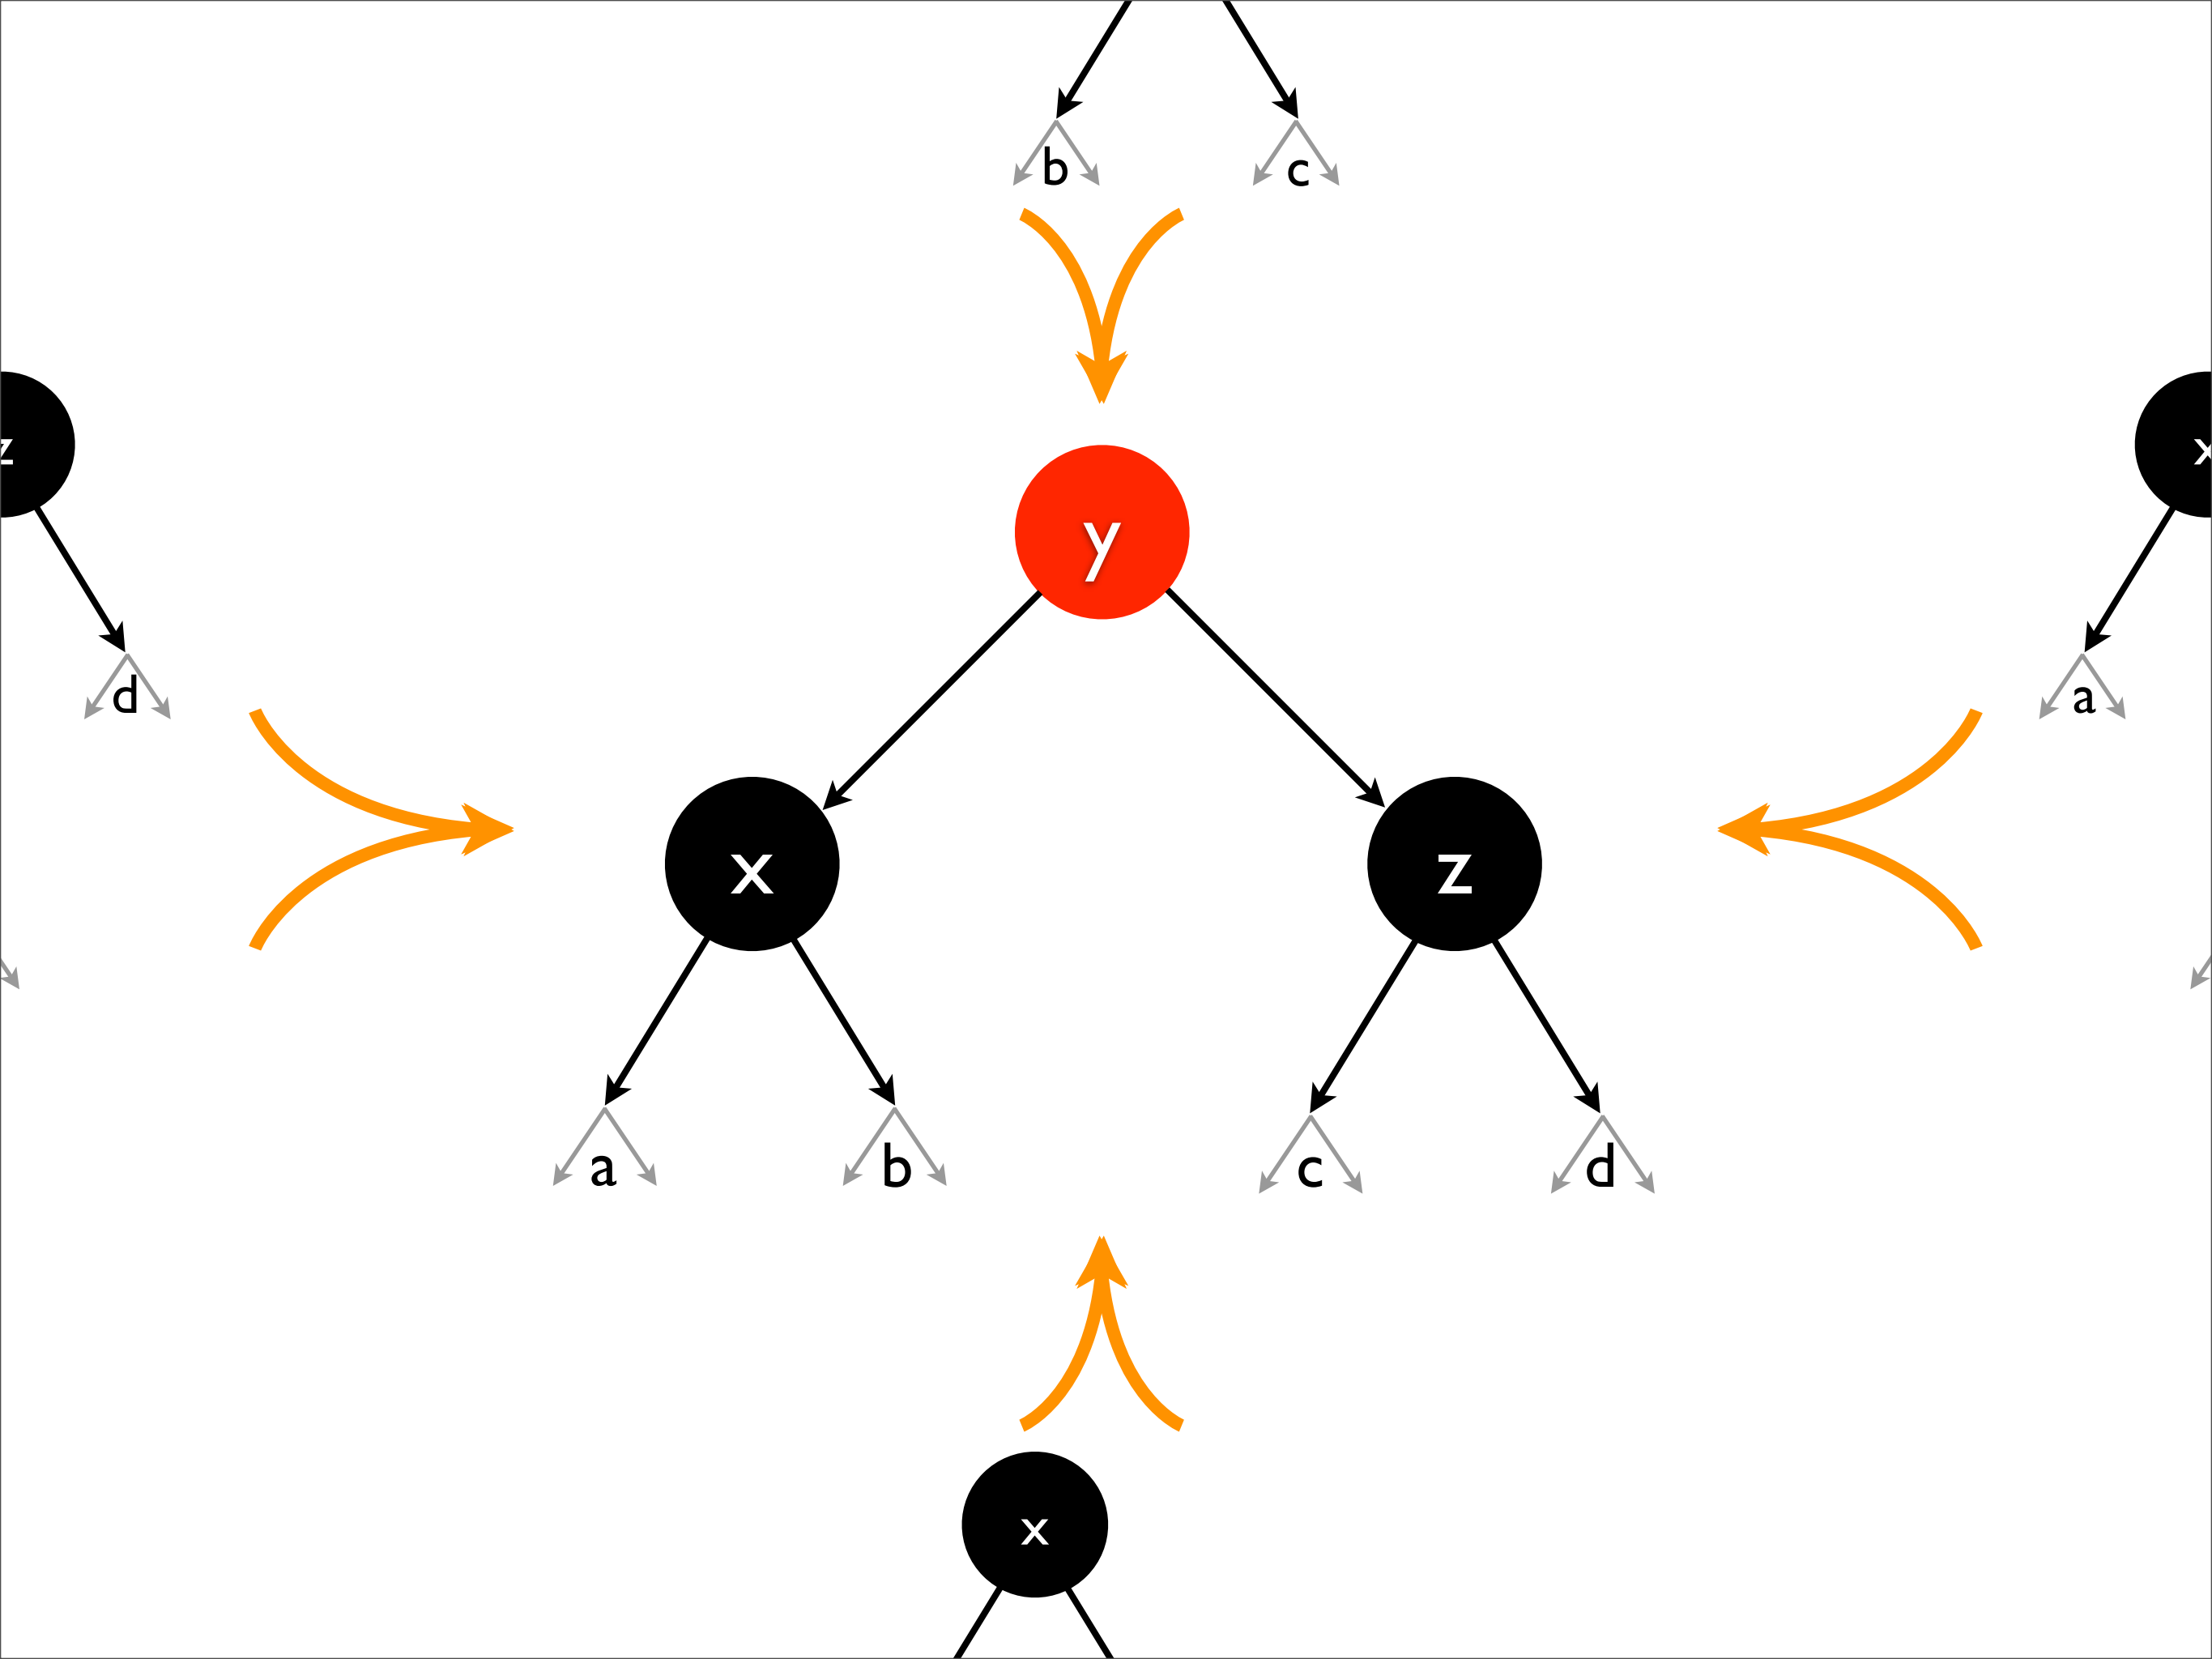
```
case class Node[K : Ordering, +V](isBlack: Boolean,  
                                   left: Tree[K, V],  
                                   key: K, value: V,  
                                   right: Tree[K, V])  
  extends Tree[K, V] {  
    ...  
}
```




```
case class Leaf[K : Ordering]()  
  extends Tree[K, Nothing] {  
    val isBlack = true  
    ...  
}
```

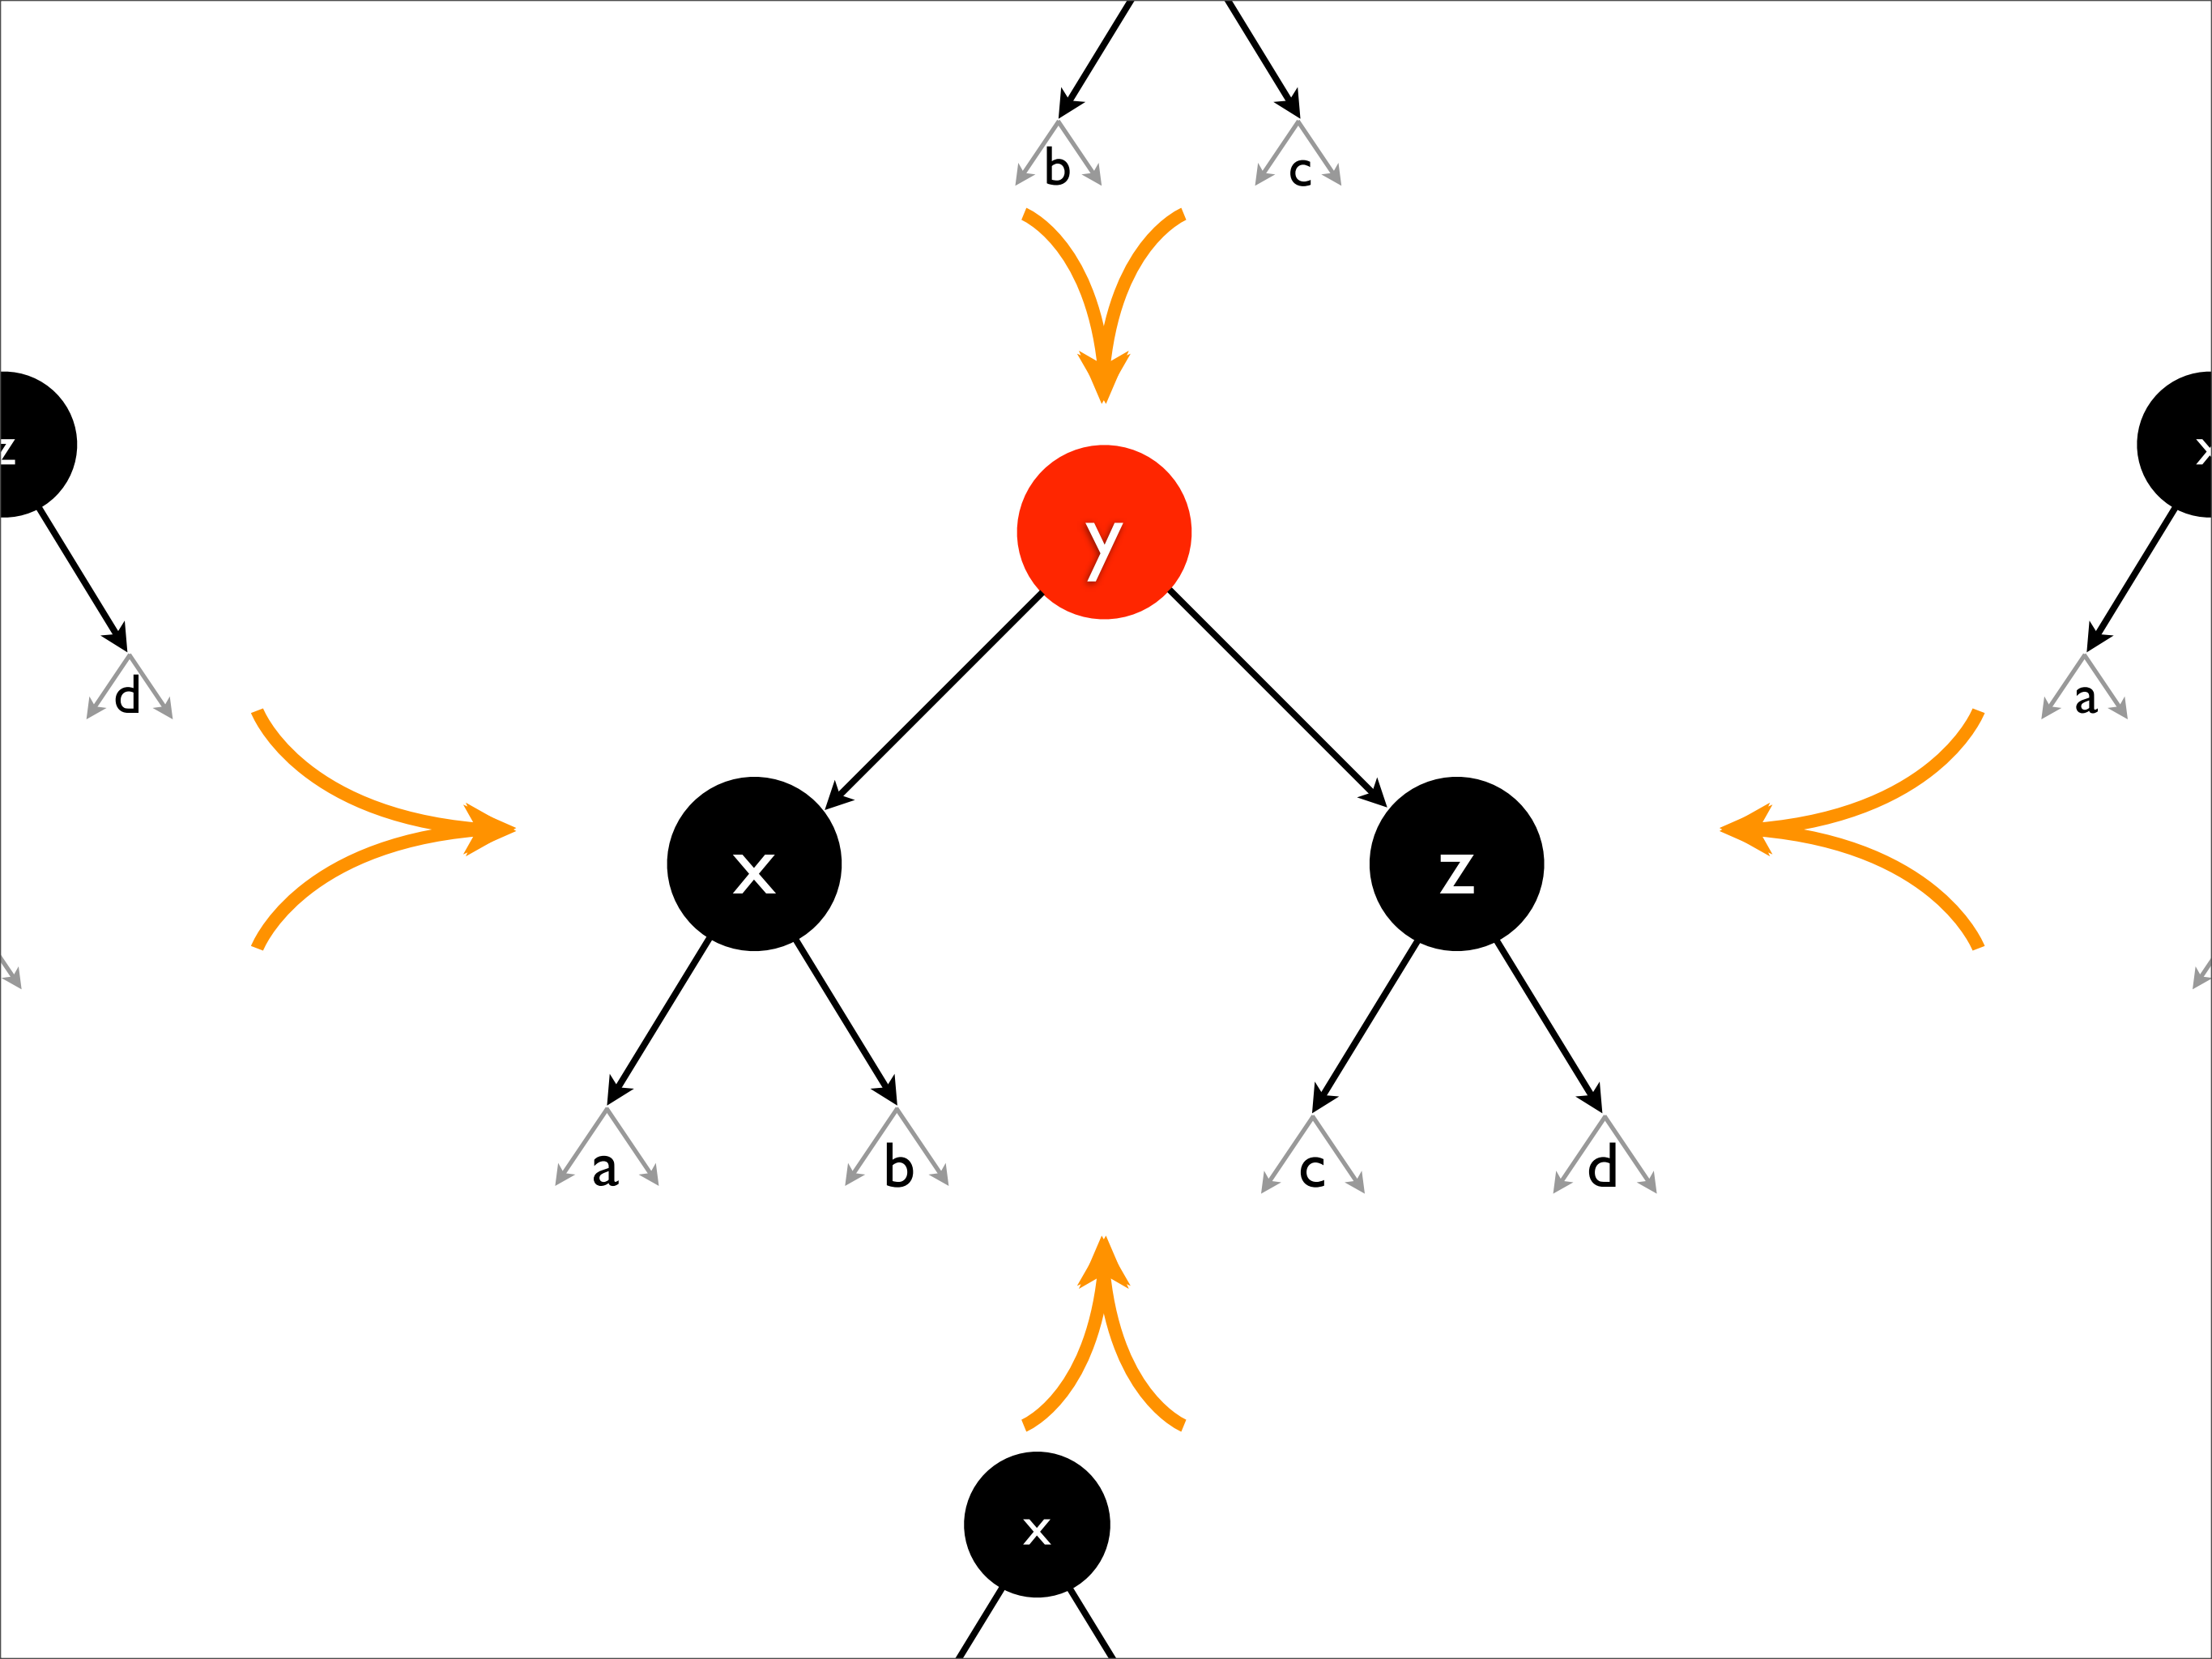





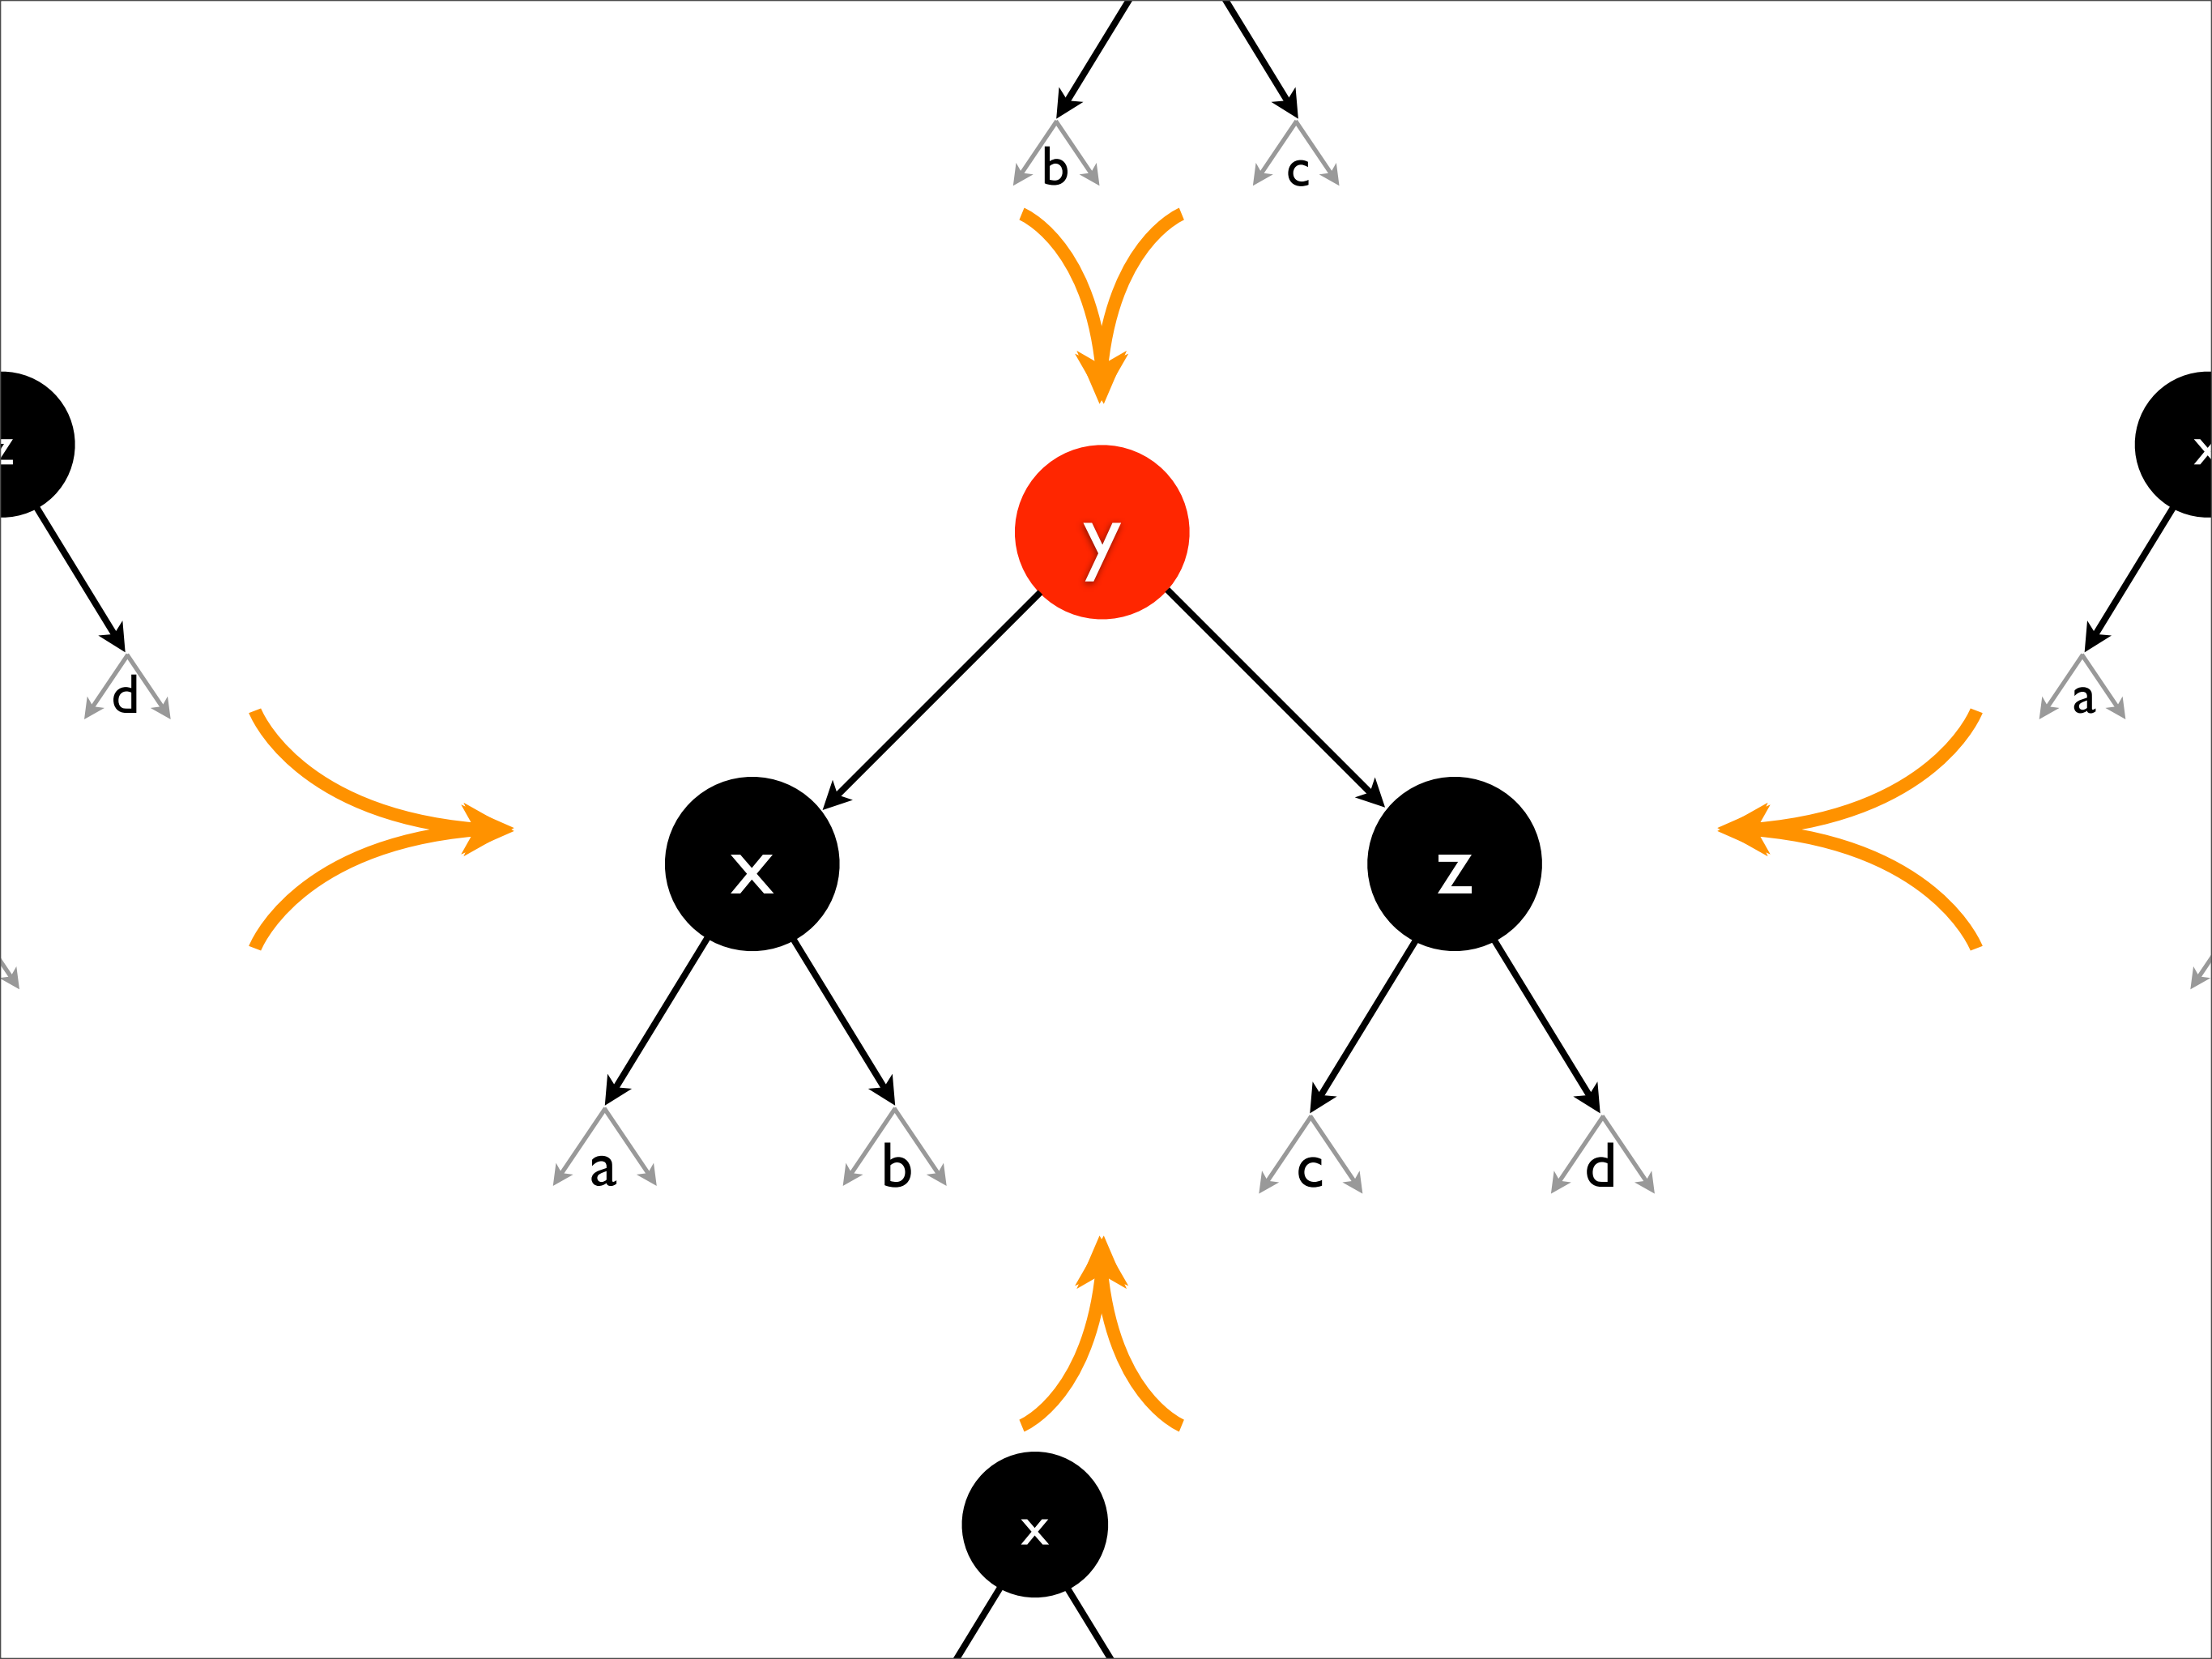





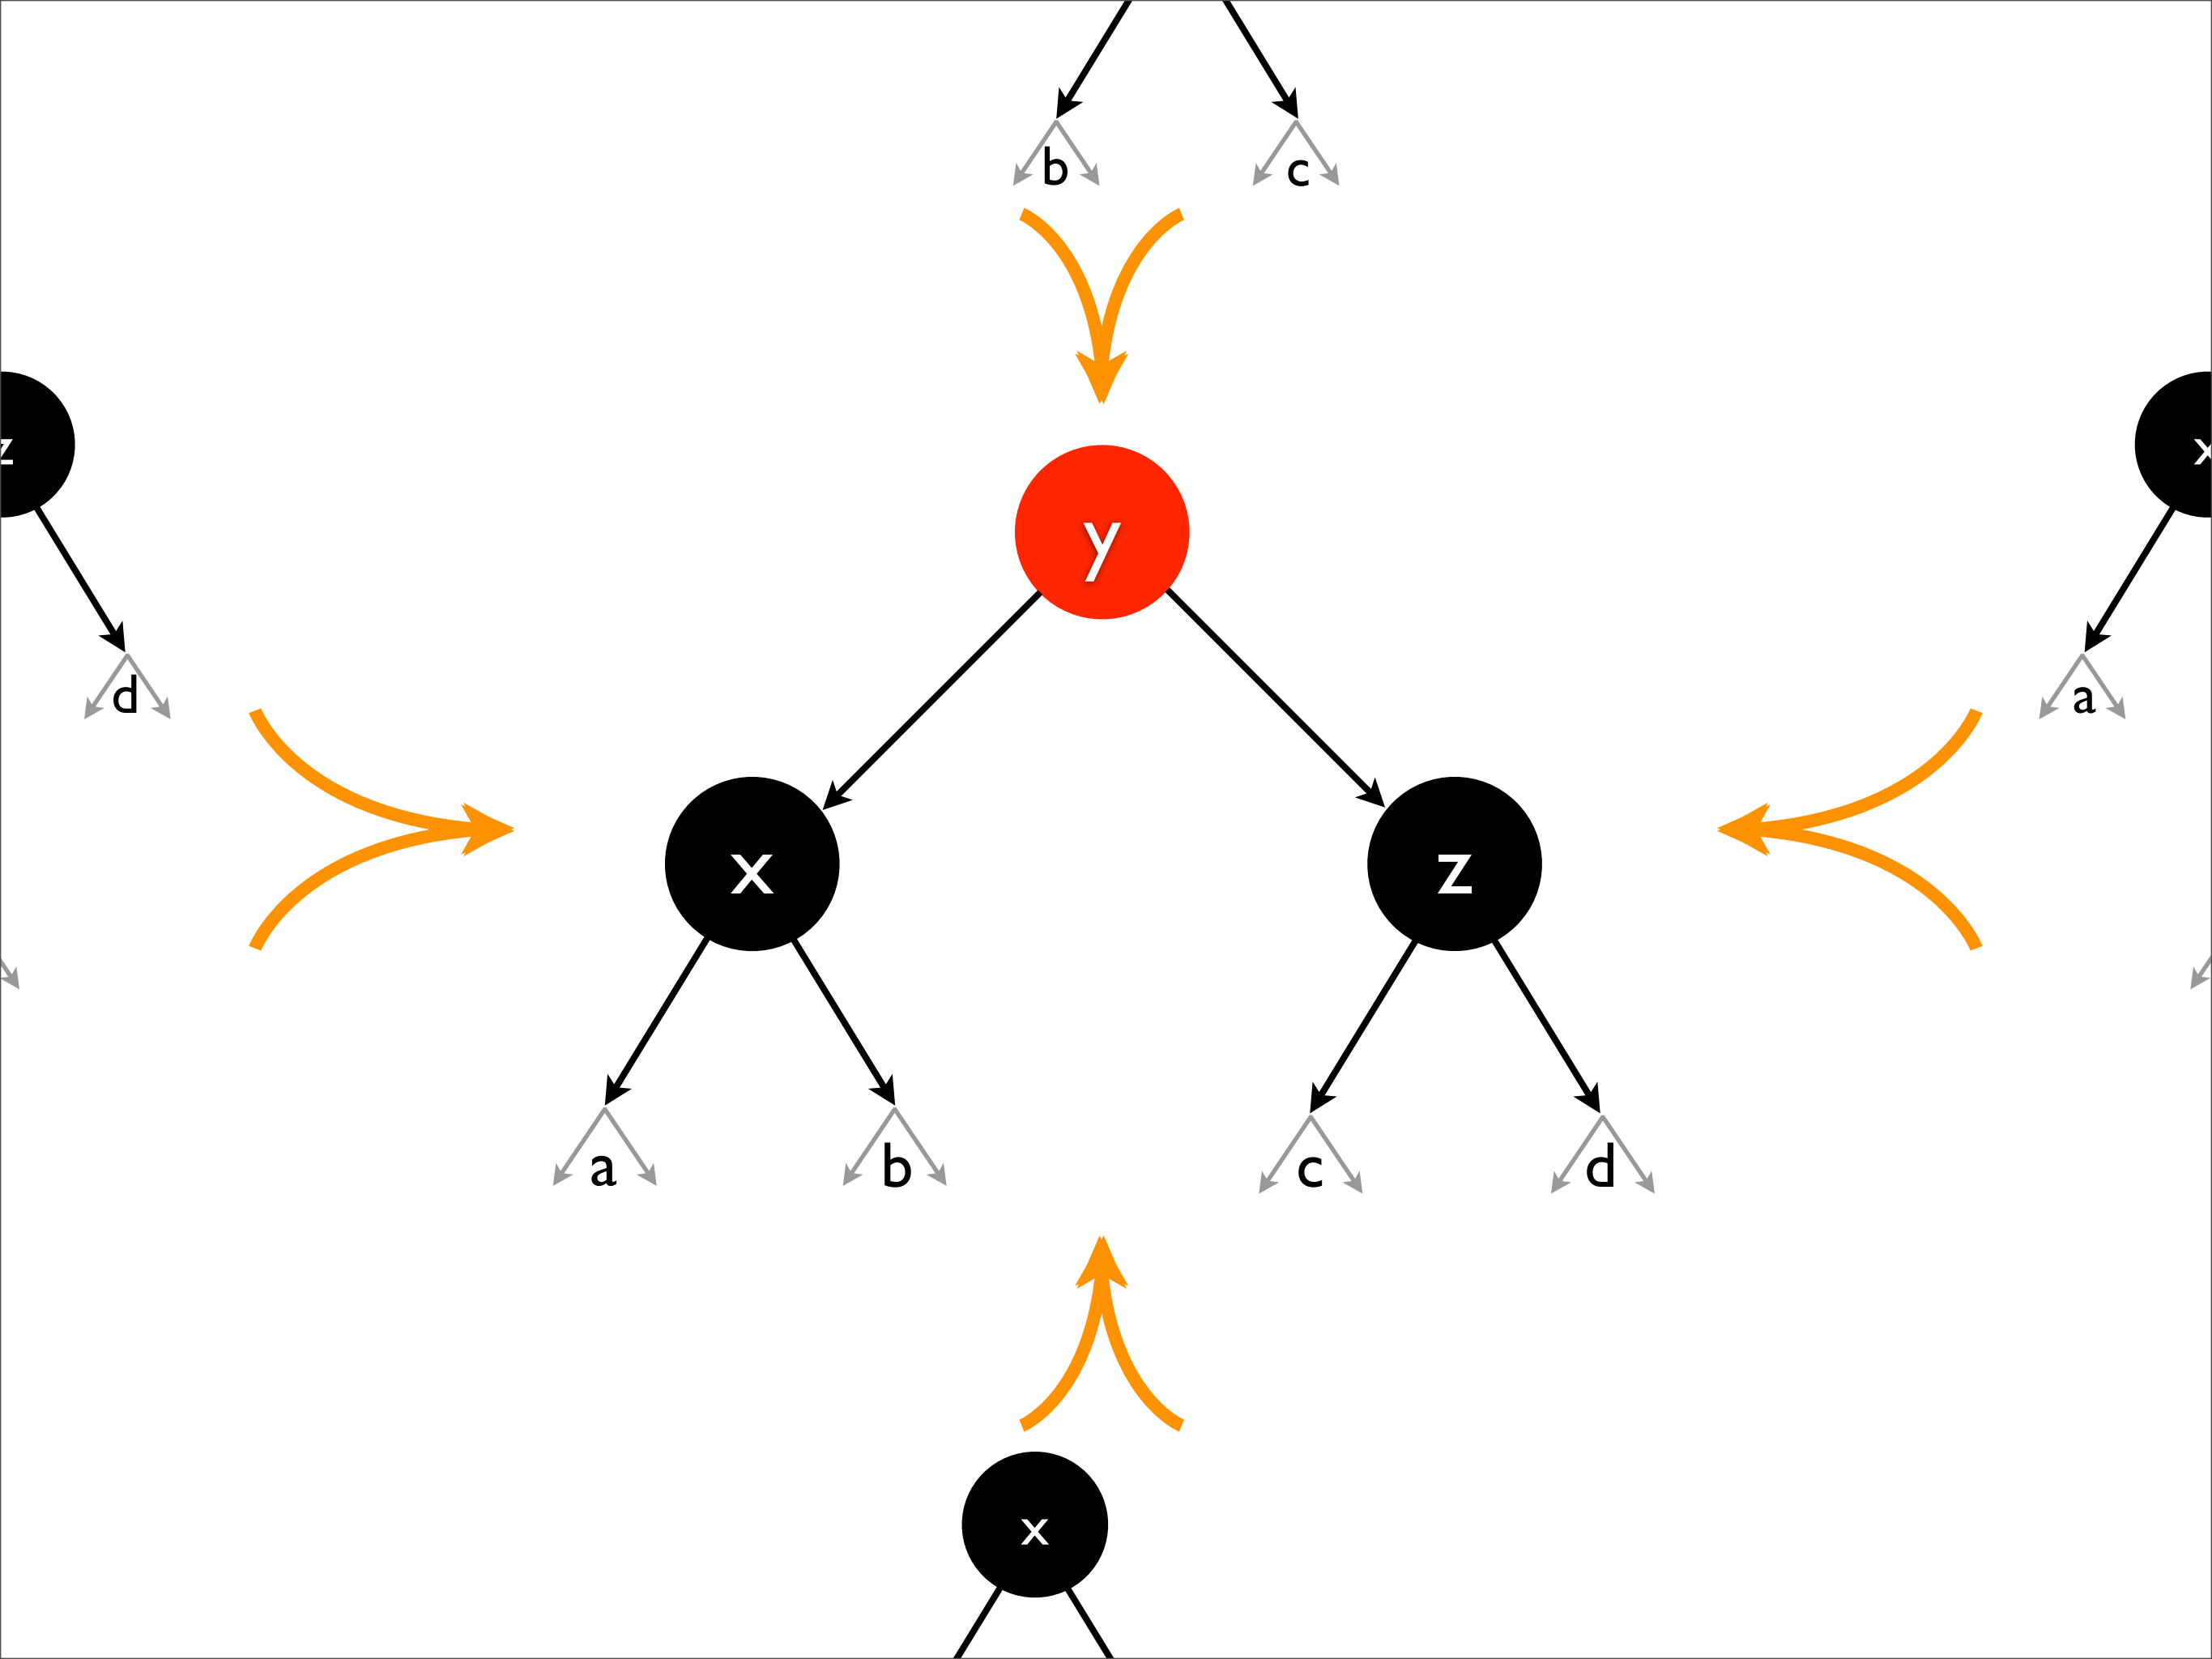












```

def balance[K : Ordering, V](isBlack: Boolean, left: Tree[K, V], key: K,
                             value: V, right: Tree[K, V]) = {
  (isBlack, left, key, value, right) match {
    case (true, Node(false, Node(false, a, xk, xv, b), yk, yv, c), zk, zv, d) =>
      Node(false, Node(true, a, xk, xv, b), yk, yv, Node(true, c, zk, zv, d))

    case (true, Node(false, a, xk, xv, Node(false, b, yk, yv, c)), zk, zv, d) =>
      Node(false, Node(true, a, xk, xv, b), yk, yv, Node(true, c, zk, zv, d))

    case (true, a, xk, xv, Node(false, Node(false, b, yk, yv, c), zk, zv, d)) =>
      Node(false, Node(true, a, xk, xv, b), yk, yv, Node(true, c, zk, zv, d))

    case (true, a, xk, xv, Node(false, b, yk, yv, Node(false, c, zk, zv, d))) =>
      Node(false, Node(true, a, xk, xv, b), yk, yv, Node(true, c, zk, zv, d))

    case (isBlack, a, xk, xv, b) => Node(isBlack, a, xk, xv, b)
  }
}

```

```

isBlack: Boolean, left: Tree[K, V], key: K,
           value: V, right: Tree[K, V]) = {
  right) match {
    Node(false, a, xk, xv, b), yk, yv, c), zk, zv, d) =>
      a, xk, xv, b), yk, yv, Node(true, c, zk, zv, d))

    , xk, xv, Node(false, b, yk, yv, c)), zk, zv, d) =>
      a, xk, xv, b), yk, yv, Node(true, c, zk, zv, d))

    de(false, Node(false, b, yk, yv, c), zk, zv, d)) =>
      a, xk, xv, b), yk, yv, Node(true, c, zk, zv, d))

    de(false, b, yk, yv, Node(false, c, zk, zv, d))) =>
      a, xk, xv, b), yk, yv, Node(true, c, zk, zv, d))

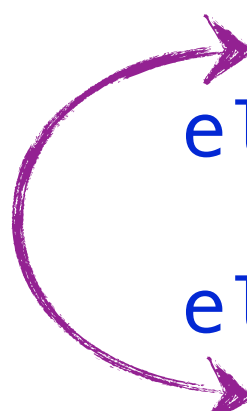
    b) => Node(isBlack, a, xk, xv, b)

```

```
case class Node[K : Ordering, +V](...) extends ... {  
  def +[A >: V](pair: (K, A)): Tree[K, V] = {  
    val (k2, v2) = pair  
  
    if (key > k2)  
      balance(isBlack, left + pair, key, value, right)  
    else if (key == k2)  
      Node(isBlack, left, k2, v2, right)  
    else  
      balance(isBlack, left, key, value, right + pair)  
  }  
}
```



```
case class Node[K : Ordering, +V](...) extends ... {  
  def +[A >: V](pair: (K, A)): Tree[K, V] = {  
    val (k2, v2) = pair  
  
    if (key > k2)  
      balance(isBlack, left + pair, key, value, right)  
    else if (key == k2)  
      Node(isBlack, left, k2, v2, right)  
    else  
      balance(isBlack, left, key, value, right + pair)  
  }  
}
```




A little of both...



Bitmapped Vector Trie

Vector(0 to 160000 by 5: _*)



Complexity

$O(1)$	$O(\log n)$	$O(n)$
append		concat
first		insert
last		prepend
nth		
update		

$$O(\log_{32} n)$$

$$O(\log_{32} n) \approx O(1)$$


Anatomy


- Start with an array with max length 32
 - Copy on write
- Array of arrays, max length 32
 - Array of array of arrays, max length 32
 - ...
- Maximum depth is 7!


a


a	b	c	d
---	---	---	---


a	b	c	d
---	---	---	---















```
sealed trait Case {  
  type Self <: Case  
  
  val shift: Int  
  
  def apply(i: Int): Array[AnyRef]  
  def updated(i: Int, obj: AnyRef): Self  
  
  def :+(node: Array[AnyRef]): Case  
  def pop: (Case, Array[AnyRef])  
}
```

```
case class One(trie: Array[AnyRef]) extends Case {
  type Self = One

  val shift = 0

  def apply(i: Int) = trie


  def updated(i: Int, obj: AnyRef) = {
    val trie2 = copy1(trie, new Array(trie.length))
    trie2(i & 0x01f) = obj
    One(trie2)
  }

  def :+(tail: Array[AnyRef]) = Two(Array(trie, tail))

  def pop = (Zero, trie)
}
```


Modern Architectures

- Locality of reference
- Caching
 - Bite-sized data chunks
- JVM considerations
 - Heap locality



References

- [1] Okasaki; *Purely Functional Data Structures*
- [2] Okasaki; *Red-Black Trees in a Functional Setting*
- [3] Hinz & Paterson; *Finger trees, a simple general-purpose data structure*



<http://github.com/djspiewak/extreme-cleverness>

